

Progressive Growing of GANs for Improved Quality, Stability, and Variation

Contents


- Introduction
- **Progressive Growing of GANs**
- Increasing Variation using **Minibatch Standard Deviation**
- **Normalization** in Generator and Discriminator
- **Multi-Scale Statistical Similarity** for Assessing GAN Results
- Experiments

Introduction

Generative Models - overview

- Autoregressive model (e.g. PixelCNN)
 - Sharp images, slow to evaluate. No latent space
- VAEs
 - Fast to train, blurry images
- GANs
 - Sharp images, low resolution, limited variation, unstable training

Generative Models - GANs



Blue, dashed line: discriminative distribution




Black, dotted line: data distribution (P_{data})

Green solid line: generative distribution (P_g)

Generative Models - GANs

- **Challenge 1**

- If not much overlap between training and generated distributions then gradients can point in random directions.



Generative Models - GANs

- **Challenge 2**

- Little variation in results



Plot of 100 GAN Generated MNIST Figures After 100 Epochs


Generative Models - GANs

- **Challenge 3**

- High resolution harder because easier to tell apart



Low Resolution



High Resolution

Generative Models - GANs

- **Challenge 4**

- High resolution requires smaller minibatches so training less stable



High Resolution

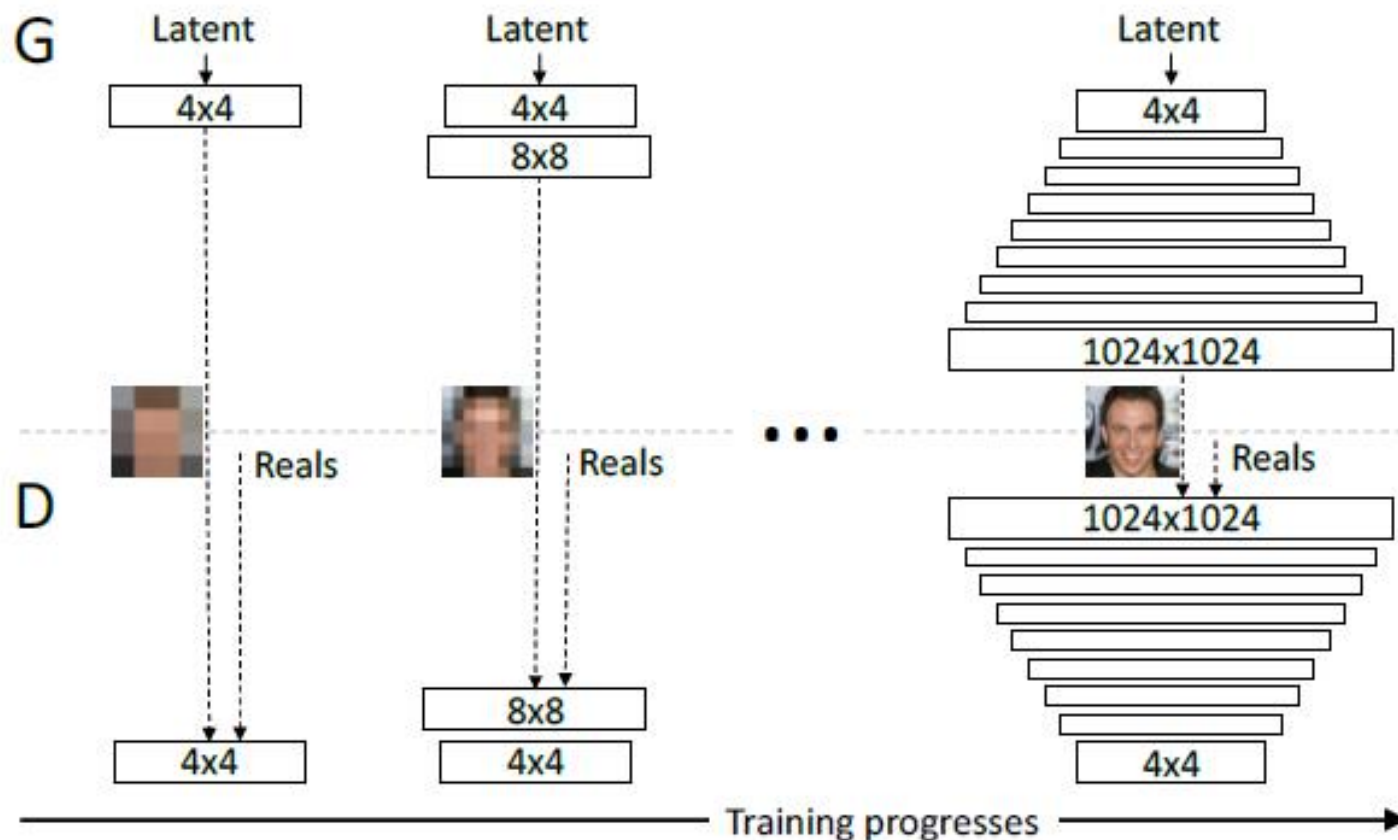
PGGAN - overview

- Key insight
 - we can **grow both the generator and discriminator progressively,**
 - starting from easier **low-resolution images,**
 - and **add new layers** that introduce higher-resolution details as the training progress

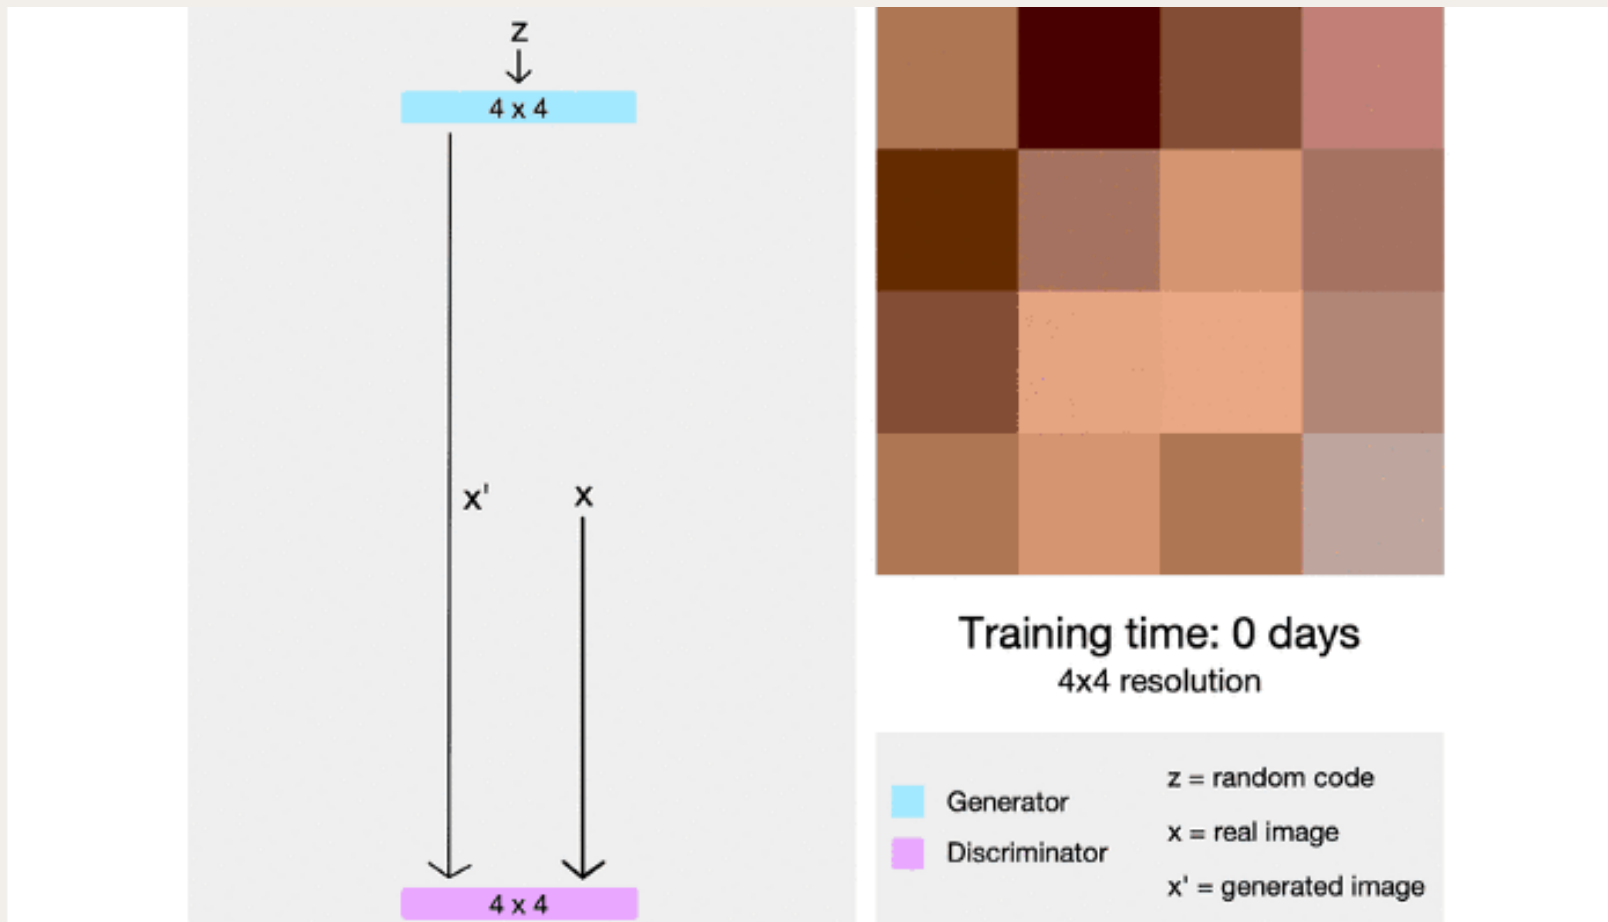


Progressive Growing of GANs

PGGAN - growing the GAN

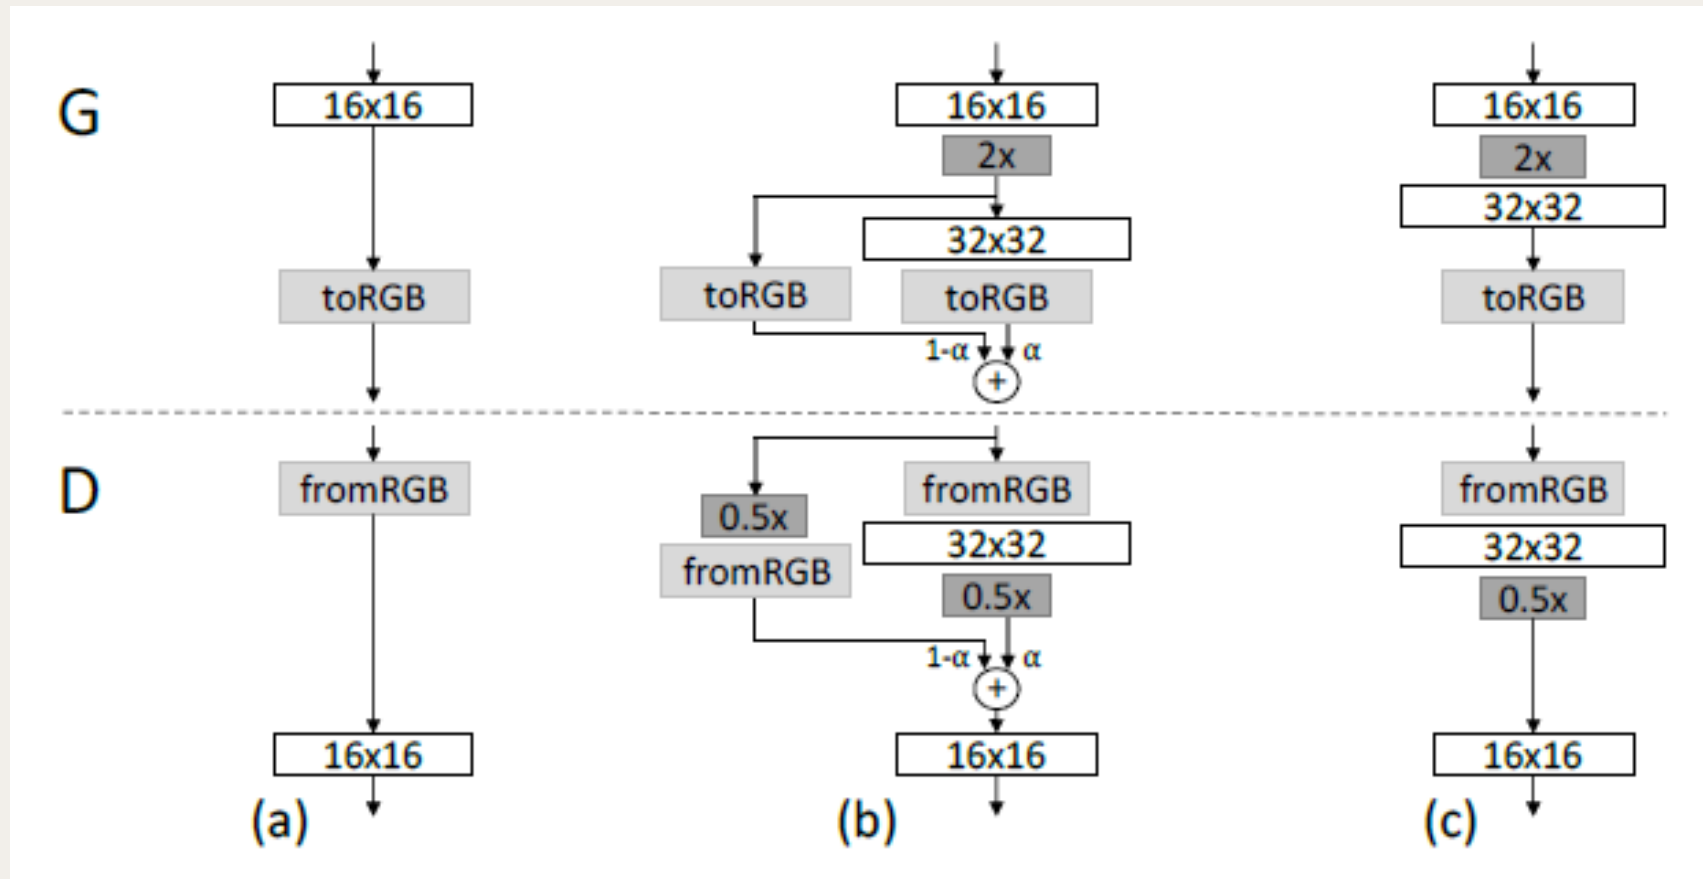


PGGAN - growing the GAN



PGGAN - fading in higher resolution layers

Fade in smoothly



Loss Function

- The author's say their work is **independent of loss function**
- Do experiments with both WGAN-GP and LSGAN

PGGAN - benefits

- Training **avoids high resolution problem** of too much divergence early on
- **Faster** training, 2-6 x faster
- Only use a **single GAN** instead of a hierarchy of GANs
- More **stable training** - more steps done at lower resolution with larger minibatches

Increasing Variation using Minibatch Standard deviation

Minibatch standard deviation

- **Minibatch discrimination** (Salimans et al. 2016)
 - Compute feature statistics **across the minibatch**
 - Encourage the minibatches of generated and training images to show similar statistics
 - Add a **minibatch layer** towards the end of the discriminator

Minibatch standard deviation

- **Minibatch standard deviation**
 - **Simplifies** the minibatch discrimination and **improves variation**
 - How to compute ?
 - Compute standard deviation for each feature in each spatial location
 - Then average over all features and spatial locations to get a single value
 - Replicate the value and concatenate it to all spatial locations and over the minibatch, yields one additional feature map

Normalization in Generator and Discriminator

Equalized Learning Rate

- Use trivial **N(0, 1)** weight initialization and **scale weights** at runtime
- $\hat{w}_i = w_i / c$
 - w_i : weights, c : per-layer normalization constant from He's initializer
- Adaptive SGD methods normalize a gradient update by its estimated standard deviation
- If some parameters have a **larger dynamic range** than others, they will take **longer to adjust**
- Our approach ensures that the **dynamic range**, and thus the **learning speed**, is the **same for all weights**


Pixelwise Feature Vector Normalization in Generator

- Normalize the feature vector in each pixel to unit length in the generator after each convolutional layer
- Prevent feature map magnitudes from getting too large


$$b_{x,y} = a_{x,y} / \sqrt{\frac{1}{N} \sum_{j=0}^{N-1} (a_{x,y}^j)^2 + \epsilon}, \text{ where } \epsilon = 10^{-8}$$

Multi-Scale Statistical Similarity for Assessing GAN Results

- **MS-SSIM**: Good at identifying global mode collapse, **not good for local mode collapse** like on colors and textures
- Do MS-SSIM on local patches drawn from **Laplacian pyramid** representations of generated and target images



Mode collapsing



Multi-Scale Statistical Similarity for Assessing GAN Results

- Sample 16,384 images and extract 128 descriptors from each level of the Laplacian pyramid. Each descriptor is a 7x7 pixel neighborhood with 3 color channels
- Compute **sliced Wasserstein distance** between samples. **Smaller distance** means that at that level of resolution training images and generator samples have **similar variation**


Experiments

Importance of Individual Contributions in Teams of Statistical Similarity

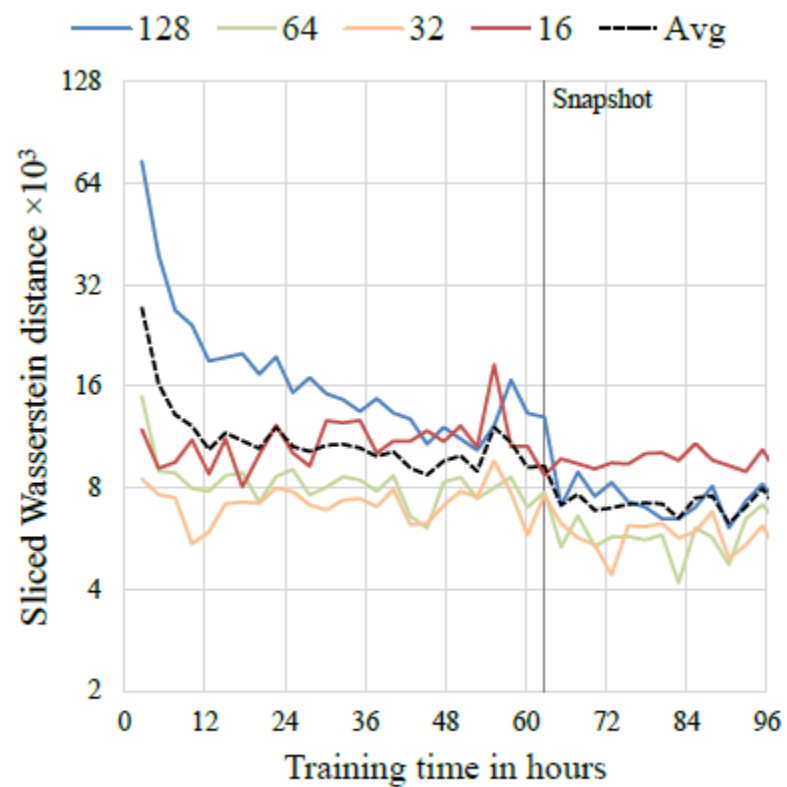
Training configuration	CELEBA						LSUN BEDROOM					
	Sliced Wasserstein distance $\times 10^3$					MS-SSIM	Sliced Wasserstein distance $\times 10^3$					MS-SSIM
	128	64	32	16	Avg		128	64	32	16	Avg	
(a) Gulrajani et al. (2017)	12.99	7.79	7.62	8.73	9.28	0.2854	11.97	10.51	8.03	14.48	11.25	0.0587
(b) + Progressive growing	4.62	2.64	3.78	6.06	4.28	0.2838	7.09	6.27	7.40	9.64	7.60	0.0615
(c) + Small minibatch	75.42	41.33	41.62	26.57	46.23	0.4065	72.73	40.16	42.75	42.46	49.52	0.1061
(d) + Revised training parameters	9.20	6.53	4.71	11.84	8.07	0.3027	7.39	5.51	3.65	9.63	6.54	0.0662
(e*) + Minibatch discrimination	10.76	6.28	6.04	16.29	9.84	0.3057	10.29	6.22	5.32	11.88	8.43	0.0648
(e) Minibatch stddev	13.94	5.67	2.82	5.71	7.04	0.2950	7.77	5.23	3.27	9.64	6.48	0.0671
(f) + Equalized learning rate	4.42	3.28	2.32	7.52	4.39	0.2902	3.61	3.32	2.71	6.44	4.02	0.0668
(g) + Pixelwise normalization	4.06	3.04	2.02	5.13	3.56	0.2845	3.89	3.05	3.24	5.87	4.01	0.0640
(h) Converged	2.42	2.17	2.24	4.99	2.96	0.2828	3.47	2.60	2.30	4.87	3.31	0.0636

Importance of Individual Contributions in Teams of Statistical Similarity

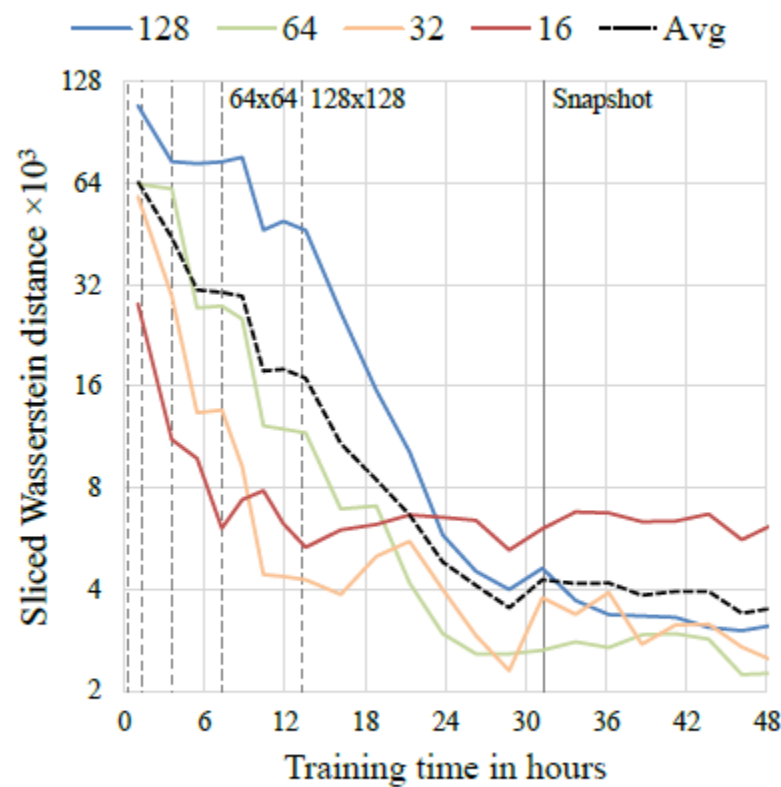
Training configuration	CELEBA						LSUN BEDROOM					
	Sliced Wasserstein distance $\times 10^3$					MS-SSIM	Sliced Wasserstein distance $\times 10^3$					MS-SSIM
	128	64	32	16	Avg		128	64	32	16	Avg	
(a) Gulrajani et al. (2017)	12.99	7.79	7.62	8.73	9.28	0.2854	11.97	10.51	8.03	14.48	11.25	0.0587
(b) + Progressive growing	4.62	2.64	3.78	6.06	4.28	0.2838	7.09	6.27	7.40	9.64	7.60	0.0615
(c) + Small minibatch	75.42	41.33	41.62	26.57	46.23	0.4065	72.73	40.16	42.75	42.46	49.52	0.1061
(d) + Revised training parameters	9.20	6.53	4.71	11.84	8.07	0.3027	7.39	5.51	3.65	9.63	6.54	0.0662
(e*) + Minibatch discrimination	10.76	6.28	6.04	16.29	9.84	0.3057	10.29	6.22	5.32	11.88	8.43	0.0648
(e) Minibatch stddev	13.94	5.67	2.82	5.71	7.04	0.2950	7.77	5.23	3.27	9.64	6.48	0.0671
(f) + Equalized learning rate	4.42	3.28	2.32	7.52	4.39	0.2902	3.61	3.32	2.71	6.44	4.02	0.0668
(g) + Pixelwise normalization	4.06	3.04	2.02	5.13	3.56	0.2845	3.89	3.05	3.24	5.87	4.01	0.0640
(h) Converged	2.42	2.17	2.24	4.99	2.96	0.2828	3.47	2.60	2.30	4.87	3.31	0.0636



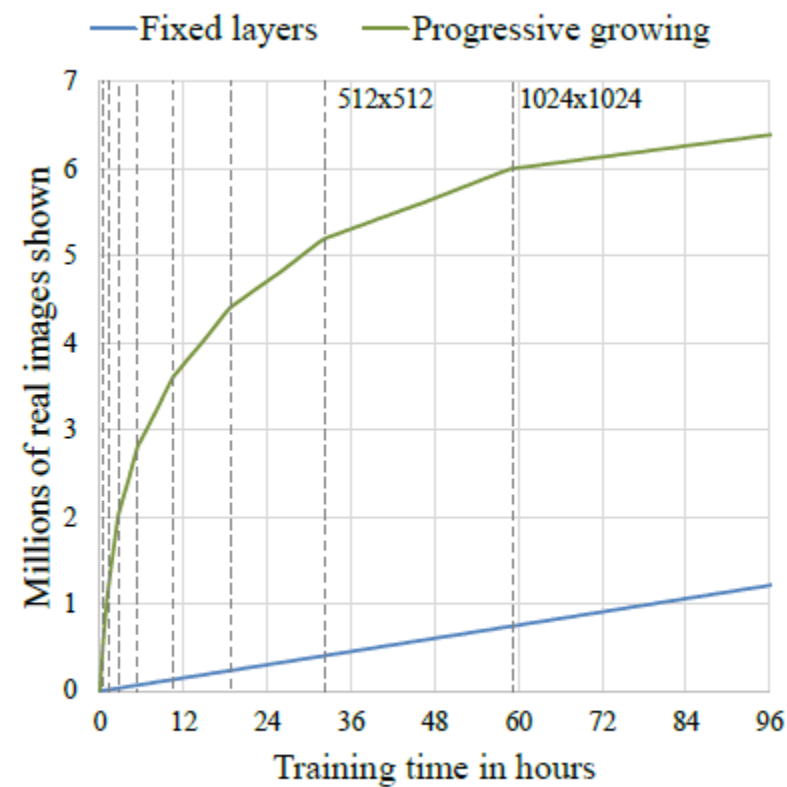
Convergence and Training Speed



(a)



(b)




(c)

High-Resolution Image Generator using CelebA-HQ Dataset



High-Resolution Image Generator using CelebA-HQ Dataset



LUSN Results



Mao et al. (2016b) (128×128)

Gulrajani et al. (2017) (128×128)

Our (256×256)

LUSN Results



Q & A